

Dictionary Learning for Disentangling Concepts in Vision Transformers.

Romeo Valentin

January 14, 2024

Contents

Abstract: Can I write an abstract here?

1 Introduction.

The Transformer architecture has been successfully deployed for a variety of applications. However, as with other machine learning methods, its reliability is not guaranteed. Reliability can for example be measured as a combination of in-distribution performance, performance under distributional shift, out-of-distribution detection and/or performance, uncertainty calibration, and robustness to adversarially chosen input perturbations.

We hypothesize that for robustness and generalizability, the model must learn a faithful representation of the world, which goes beyond mere statistical inference.

Humans similarly have a model of the world, and although it is generally not easy to write down, we have developed natural language as a way to communicate information between each other. It is therefore natural to wonder whether the models potentially emerging in large Transformer networks are or can be aligned with a human model of the world.

Let us consider the special case of Vision Transformers. Typically, models are trained to produce a sequence of high-dimensional embeddings, which can serve as inputs to downstream tasks like classification. It is therefore clear that these embeddings have a variety of information stored in them.

The goal is now try to “disentangle” these embeddings in order to provide a interpretable/human-aligned/monosemantic view of the embeddings, which stay faithful to the original content of the embedding (but might sacrifice some of the detailed information). This goal of “disentanglement” from

a “block of data” goes back to Independent Component Analysis (ICA), and was picked up again in the context of machine learning (see e.g. beta-VAE (Higgins et al. 2017)). “Disentanglement” here is defined as having a set of components that are both conditionally independent and align with a notion of human interpretability. However, an important result in 2020 proved that without further assumptions there is an infinite set of components which satisfy the conditional independence property, so requiring this alone is not enough to get interpretable components for free (Locatello et al. 2019).

However, in this work we propose two extra conditions with which we aim to overcome the theoretical concerns:

1. small number of “active” concepts, and
2. encodable using natural language.

In the next sections we first discuss some theoretical details for both conditions, and then describe how we actually implement each condition, i.e through dictionary learning/the k-SVD algorithm (Section ??) and token assignment (Section ??).

1.1 Small number of “active” concepts.

The first condition is motivated easily. It turns out to be easy to project a dataset onto a new basis of orthogonal vectors, i.e. “statistically independent concepts”, so that each datapoint is a sum of these orthogonal vectors – this is indeed just the singular value decomposition (SVD), i.e. $Y = U\Sigma V^\top$ where each column of Y is a datapoint, and each column of U is a “concept vector”.

However, in this case each datapoint is constructed of an “infinite sum” of concept vectors, whereas we believe that any data sample only has a relatively small number of active concepts. Further, the SVD is designed to find something like “Principal Components”, which in our language would translate to abstract concepts that are represented in the majority of datapoints. However, we want neither of these. Instead, we would like to represent each datapoint y for example as a sum $\sum_{j \in \mathcal{J}} \vec{d}_j \cdot \lambda_j$ where $|\mathcal{J}|$ is small. In other words, we try to write our data matrix $Y = [y_1 \dots y_{|Y|}]$ as a product of dictionary vectors $D = [d_1 \dots d_{|D|}]$ and a sparse assignment matrix X such that $Y \approx DX$.¹

¹Note that having a linear dependence of Y wrt D is a design choice, and one may also consider more nonlinear relations. However, for tractability we restrict ourselves to the linear setting.

It turns out that this problem has been studied previously, and is known as “Dictionary Learning”. A popular algorithm for dictionary learning is the “K-SVD” algorithm (Aharon, Elad, and Bruckstein 2006), which essentially alternates between updating the set of dictionary vectors and updating the assignment matrix. We will go into more details in Section ??.

Recently, it has also been proposed to use a Sparse Autoencoder, e.g. as presented in (Bricken et al. 2023). We briefly compare the approaches below.

1.2 Natural language encoding.

- The additional sparsity constraint is enough to overcome the theoretical problems brought up by Locatello et al. (2019).
- However, it remains unclear whether this is enough for extracting dictionary or concept vectors with the kind of human-aligned interpretability that we are looking for.
- Indeed, even if they were sufficiently aligned, it is not clear how we can map the concept vectors to natural language representations, or how we can test whether a language-to-embedding mapping is “faithful” or correct in some sense.
- However, we believe that the recently gained availability to large language models (LLMs) gives us a way to overcome these problems.

Proposition 1: *Consider the setup presented in Fig. ??, and consider that the disentangling method works well in the sense described in the previous section. Then, we consider a token mapping “faithful” if the diagram in Fig. ?? commutes.*

Being able to “score” a given token mapping now may make it possible to actually construct the token mapping. We go into more detail on the approaches in Section ??.

2 Dictionary Learning with the K-SVD algorithm.

In this section we first discuss some theoretical properties and considerations of the K-SVD algorithm, and then discuss the contributions we have made and published in the KSVD.jl package. Then we will briefly discuss the K-SVD algorithm versus using a Sparse Autoencoder.




Figure 1: Commutative diagram showing the flow of semantic information. Given captioned images, a ViT, an LLM, and a disentangling method, we can therefore create a token mapping.

2.1 The algorithm, and theoretical considerations.

The K-SVD algorithm solves the problem of dictionary learning outlined above, namely finding a matrix decomposition $Y \approx DX$ where $\text{size}(D, 2) > \text{size}(D, 1)$ (i.e. an “overcomplete” dictionary), and X is sparse. It can be understood as a generalization to the k-means algorithm, but allowing any datapoint to be associated with *multiple centroids*.

Typically, iterative dictionary learning algorithms alternate between two steps:

Sparse coding: Given a fixed dictionary, for each data sample find a small subset of dictionary vectors and factors such that $y_i \approx \sum_{j \in \mathcal{J}} x_j \cdot d_j$. Note that this is a non-convex problem (and indeed NP-hard), and typically heuristics like greedy search are used. Typically used algorithms are (Orthogonal) Matching Pursuit, Basis Pursuit, or FOCUSS. In our implementation, we use Matching Pursuit.

Dictionary update: Given a coding matrix X , we can again update the dictionary elements. Several approaches exist – for example, a simple gradient descent based approach can be used to optimize $\min_D \|Y - DX\|$, where X stays fixed. However, Aharon, Elad, and Bruckstein (2006) propose updating each dictionary element independently by considering all the data points that are “using” a given dictionary element and then replacing the dictionary element with the dominant singular vector of these datapoints. Note that this simultaneously updates the values in X , which most other algorithms do not.

This algorithm has exceptionally few hyperparameters – indeed, it is sufficient to choose a number of dictionary elements, and a cutoff threshold

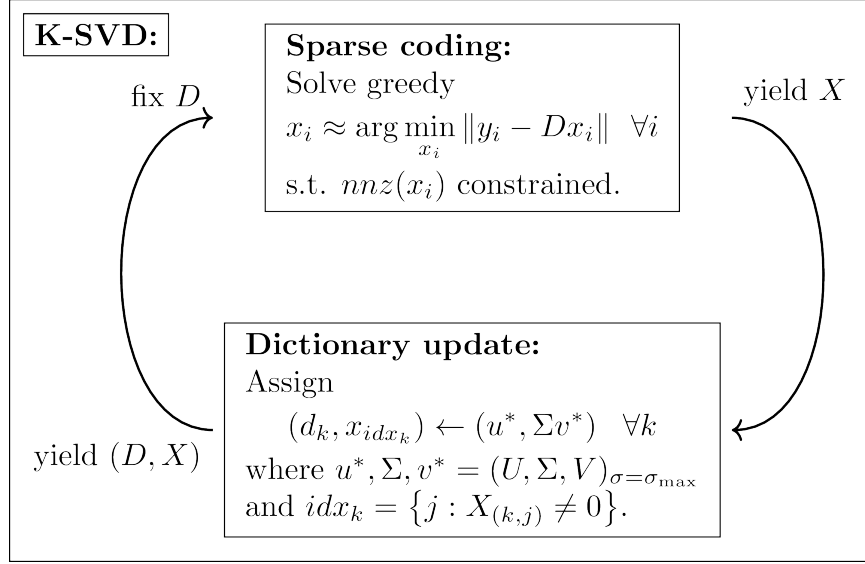


Figure 2: An overview of the K-SVD algorithm.

for sparse coding.² A mere two hyperparameters! Further, utilizing the (truncated) svd algorithm typically leads to fast convergence and good computational efficiency, since the svd is theoretically “optimal” in some sense, and heavily optimized.

A note on convergence. It is noteworthy that K-SVD has a convergence guarantee *only to a local minimum*, and only if the sparse coding step *is solved optimally*. However, as is somewhat typical for machine learning, we have some hope that for large enough datasets the loss function becomes practically convex, so that the algorithm converges to a global minimum.

2.2 KSVD.jl : A highly optimized K-SVD implementation.

Note: While we believe the timings described in this section to be roughly correct, we have to run additional tests to make sure we also achieve the same convergence.

Bricken et al. (2023) mention the K-SVD algorithm, however deem it computationally infeasible to apply to large datasets with millions (or billions) of samples. And indeed, current implementations seem not up to the

²Alternatively, we can specify a maximum number of dictionary elements per data sample.

task; the implementation available as `sklearn.decompositions.MinibatchDictionaryLearning`, which extensively leverages `numpy` and `joblib`, takes over 3 minutes for ten iterations on a dataset with ten thousand elements (despite multi-threading).

For this reason, we present `KSVD.jl`, an implementation of the K-SVD algorithm in the Julia programming language (Bezanson et al. 2017). This implementation outperforms `sklearn`’s implementation by about $50\times$ when computing the same problem, can gain an additional $2\times$ by reducing the precision from `Float64` to `Float32`, and can be scaled across many compute nodes with almost linear speedup improvements³. That means if, for example, eight compute nodes are available, we can expect a speedup of $(50 \cdot 2 \cdot 8)\times = 800\times$ for moderate to large datasets. Further, `KSVD.jl` also employs several algorithmic modifications that, to the author’s knowledge, lead to faster convergence given the same number of compute operations.⁴

This speedup has been achieved through extensive benchmarking and optimization of the code, including

- careful adjustments to the execution order and small algorithmic adjustments,
- single-core optimizations like aggressive buffer preallocations, exploiting cache locality, improving the memory layout, and reducing memory movements,
- careful multi-threading using small batch updates with frequent cross-communication implemented with Julia’s efficient task scheduling,
- a custom multi-threaded dense-sparse matrix multiplication implementation (`ThreadedDenseSparseMul.jl`),
- pipelined GPU-offloading for large matrix multiplications (currently unused in the fastest version),
- a custom distributed executor allowing to spread the computation over many compute nodes.

Over-the-thumb estimation of computational requirements for large dataset. To illustrate the (theoretical) execution times on a large dataset, let us estimate the time to compute 10 and 100 iterations of the K-SVD algorithm on embeddings from the OpenCLIP dataset, which has 400 million

³Unlike the previous two numbers, this hasn’t been properly tested yet.

⁴A more detailed study on this is to follow.

samples. We will consider having a cluster with 8 nodes and 64 cores each, and compute in `Float32` precision.

As a datapoint, let’s consider measurements from my 16 thread (8 core) Intel i7 mobile processor, which achieves 10 iterations on 800’000 samples in about 120 seconds. Using the setup above, we have about 32 times more compute resources and a 500 times larger problem, which yields an estimated runtime of $120\text{sec} \cdot \frac{500}{32} = 1875\text{sec} \approx 0.5\text{h}$, and similarly 5h for 100 iterations etc.

Dictionary Learning vs Sparse Autoencoder.

Method	k-svd	gradient
find sparse assignments via. . .	greedy search	L_1 loss
number of hyperparameters	1/2	many
“one” way to do it	yes	no
can get stuck in local minima?	yes (?)	yes (?)
runtime?	Comparison tbd. . .	
quality of result?	Comparison tbd. . .	

3 Creating the token mapping.

In the previous section, we saw how we can extract “concept” or “dictionary” vectors that explain the embeddings generated by a ViT. Crucially, we assume for now that each dictionary vector can be mapped to a human concept. This is a strong hypothesis, as it requires (i) that the introduction of the sparsity constraint overcomes the concerns mentioned in Locatello et al. (2019), and that (ii) the k-SVD algorithm does a sufficiently good job at solving the problem.

However, even in this case it is not trivial to actually find a way to assign the human concept. We propose several ways how this can be achieved nonetheless. All ways rely on having *captioned images*, i.e. dataset of pairs connecting a natural-language description and an associated image. In particular, we work with the *sbucaptions* dataset (Ordonez, Kulkarni, and Berg 2011), which contains about 800 thousand captioned images, see e.g. <https://vislang.ai/sbu-explorer>. Other datasets include coco (600k), conceptual captions (3M and 12M), laion datasets (400M and 5B), and others.

The approach for all the following methods is always similar:

1. *Only* the images are fed to a ViT model, and a single embedding vector

per image is constructed (usually 768 or 1536 dimensional), yielding an embedding matrix like $Y \in \mathbb{R}^{768 \times n}$.

2. Dictionary learning (e.g. with the k-SVD algorithm) is used to rewrite Y as $Y \approx DX$ with X sparse and D overcomplete.
3. The captions are processed somehow, e.g. by computing an embedding or checking them for certain concept nouns. The result constitutes a kind of label.
4. These labels are used to assign concepts to each dictionary element in D .

Notice that these steps roughly model the diagram in Fig. ??.

We will now go over some concrete examples over how this can be achieved. We denote approaches that have been implemented in some form of proof-of-concept, and which have not.

3.1 Correlating word occurrences with sparse assignments.

We first present an approach not reliant on LLMs at all, and instead works by correlating X with word occurrences in the captions. We start by considering a list of English words, e.g. nouns from here. For each caption, we then check if the noun is contained in the caption, yielding a occurrence matrix $O \in \{0, 1\}^{n_{\text{words}} \times n_{\text{captions}}}$.

Recall now that we have $X \in \{0, 1\}^{n_{\text{dictionary}} \times n_{\text{words}}}$ and let $X_{01} \in \{0, 1\}^{n_{\text{dictionary}} \times n_{\text{captions}}}$ be the “indicator” matrix of X , i.e. $X_{01}[i, j] := (X[i, j] \neq 0)$. Then, we can compute the correlation between O and X_{01} ; informally, if a certain dictionary element has been selected for a set of images, and all the images’ captions contain the word “dog”, then the dictionary concept is at least statistically related to the concept “dog”.

More precisely, let $C = \text{cor}(O, X_{01}) \in \mathbb{R}^{n_{\text{words}} \times n_{\text{dictionary}}}$ be this correlation, reduced along the images/captions. Then, for each dictionary element we can assign e.g. the top-1 or top-3 most correlating words by simply finding the maximum correlation column-wise.

Figure 3: Finding word tokens for dictionary elements by correlating X with word occurrences in image captions.

```
(D, X) = deserialize("sbu_captions_embeddings_ksvd.jls")
X_01 = (X .!= 0)
```



```

words = readlines("english-nouns.txt")
captions = map(lowercase, readlines("captions.txt"))
O = [occursin(word, caption) for word in words,
      caption in captions]
n = size(X_01, 2); @assert size(X_01, 2) == size(O, 2)

C = cor(O, X_01; dims=2)

word_indices_per_dict_element = [
    sortperm(col; rev=true)[1:3] # indices with max correlation
    for col in eachcol(C)
]
# e.g. show words for first dictionary element
@info words[word_indices_per_dict_element[1]]

[ Info: ["boat", "ship", "sail"]

```

Using this strategy also gives us a natural way to plot the relations we have found. Let us remove all correlations in C that are not in the top-3 column-wise. Then, we immediately have an adjacency matrix of a bipartite graph, where the one set of nodes are the words, and the other are the dictionary elements. We can now use classical graph layouting techniques to draw the resulting connections, e.g. through modeling connections as forces, or using spectral approaches.

In the future we would like to pull up images associated with a dictionary vector.

Limitations:

- currently no proper scoring method
- correlations often low (< 0.1)
- currently “air” and “hair” etc are assigned to the same caption.
- currently only nouns, problem with verbs/adjectives is that they are commonly conjugated.
- not clear whether

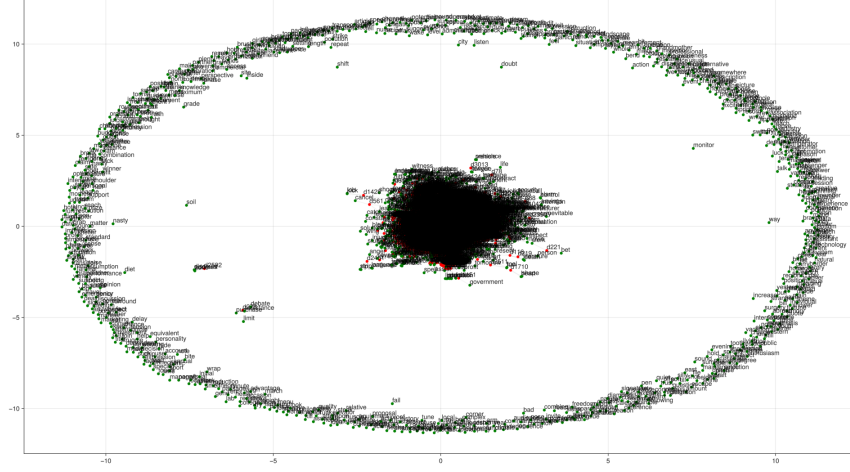


Figure 4: The resulting graph with a spring layout. The words around the outside are all words that don't have any edges.

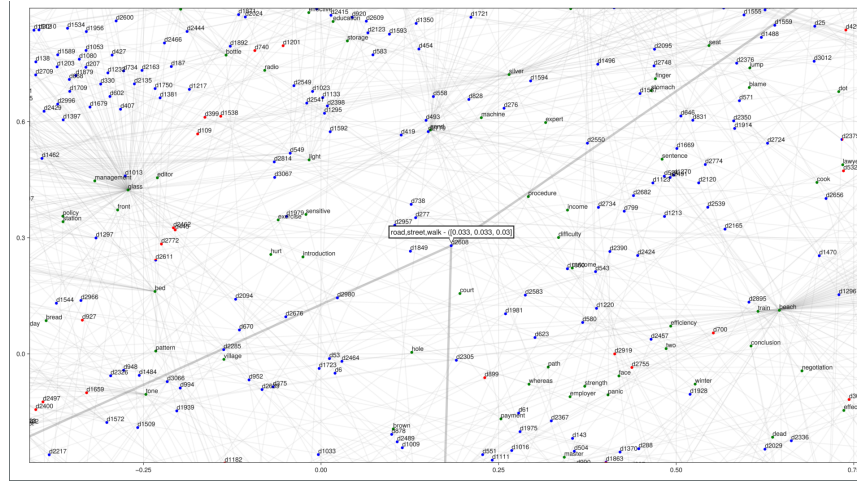


Figure 5: A zoomed in view of the figure above. A single dictionary element is associated with the words “road”, “street”, and “walk”, with correlations ≈ 0.033 each.

3.2 Minimizing semantic distance between image caption and list of descriptive tokens.

The previous approach was free of any deep learning. Generating labels through simply checking $\text{word} \in \text{caption}$ can be very noisy, both due to not modeling any grammar (e.g. declined nouns, etc), and not capturing any semantics.

However, the recent emergence of NLP modeling gives us more tools to deal with the semantic meaning of the captions. In particular, we can try to generate a “bullet point description” from the ViT embeddings and then minimize the semantic distance of this description to the caption in an embedded space.

More precisely, consider for now the case that each dictionary element has an assigned word token. Further, consider a model that computes an embedding for any text input (a relatively common type of LLM). This gives us a simple way to score any mapping from dictionary elements to tokens, and the token assignment problem reduces to a search problem.

We propose two ways to deal with this search problem – one directly trying to find tokens in the token embedding space, and one using greedy search.

Finding tokens by optimizing in the token embedding space. Let us denote the text encoding model as textenc , which takes a sequence of token embeddings, and produces an embedding, i.e. $\text{textenc} : \text{Sequence}(\mathbb{R}^e) \rightarrow \mathbb{R}^e$. The token embeddings are typically stored as a lookup table with random initialization (although they can also be the result of optimization).

We can now try to assign tokens for each dictionary element as follows:

1. We create a new matrix of randomized dictionary embeddings, i.e. a random vector for each dictionary element.
2. Then, we create a placeholder prompt, e.g.

```
An image of an object characterized by:  
- [MASK]  
- [MASK]  
- [MASK]  
...
```

This prompt is encoded into a sequence of embeddings (as is typical).

3. Now we consider an input image-caption tuple, and the corresponding sparse dictionary assignments for the image given by X . Using these assignments, we pick the corresponding embeddings from the matrix from step 1, and insert them into the prompt embedding sequence in step 2, essentially replacing the mask tokens by our own token.
4. Then we process the new sequence of token embeddings with textenc, and compare the result with calling textenc on the caption directly. We can compute e.g. a MSE score, and use the gradient to update the active token embeddings from step 1.

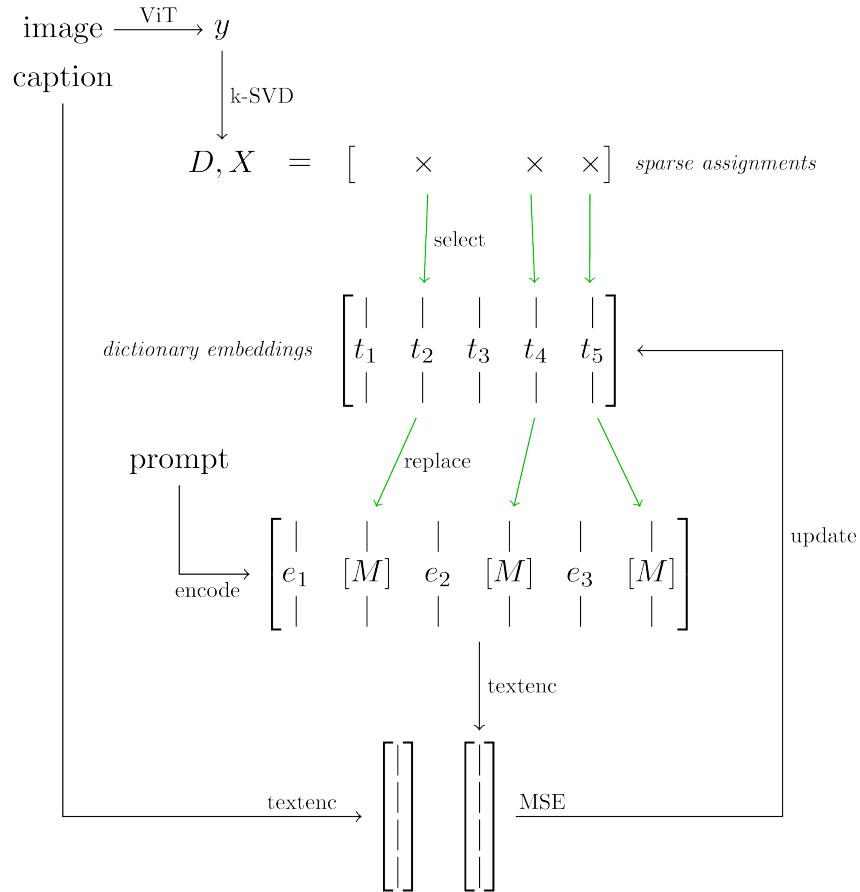


Figure 6: An overview over the token optimization in embedding space.

The goal here is that we can use the gradient to directly optimize the

dictionary embeddings in the token embedding space. Upon convergence, we would extract the natural language token from the dictionary embeddings by finding the nearest neighbor token embedding and using that token.

However, it turns out that in our experiments this approach didn’t seem to converge at all. We believe this is because models are not built to be “robust” to changes in the token embeddings at all, since these are typically fixed. Therefore, the gradient information for the dictionary embeddings might be very poorly behaved.

For this reason we turn to a more basic approach in the next section.

Finding token embeddings through greedy search. In this setup we adapt a similar approach to the previous section, however we do not rely on the gradient for optimization. Instead, we simply try to directly set the dictionary embeddings to token embeddings.

The algorithm works as follows:

1. Start by reducing the sparse assignments only to the “maximum” assignment, and use a prompt with only one mask token.
2. Then, for every image-caption-sparse-assignments triplet, try out every token the tokenizer supports, and use its embedding to replace the mask token’s embedding.
3. Keep the token that minimizes the mean-squared error (MSE) between the caption encoding and the modified prompt encoding computed over the entire dataset.
4. Add a second sparse-assignment and mask token, and repeat as above. Continue until every sparse assignment is processed.

We have had some success with this, but not processed properly yet.

3.3 Dictionary learning on a combined (word-embedding \cup image-embedding) dataset.

3.4 Finding shared dictionary vectors between word and image embeddings.

4 Current limitations.

The current setup assumes that the sparsity constraint is enough to produce dictionary vectors that only need to be aligned downstream. However, it’s

not clear that this is true. Indeed, it’s quite likely that performance would be improved if the token mapping could somehow be included in the disentangling step. This could perhaps be done by propagating a gradient back from the LLM setup and using it in the dictionary learning, which would perhaps integrate better with the Sparse Autoencoder. However, this is currently out of scope, but would certainly make a good next project.

5 Bibliography.

- Aharon, M., M. Elad, and A. Bruckstein. 2006. “K-SVD: An Algorithm for Designing Overcomplete Dictionaries for Sparse Representation.” *Ieee Transactions on Signal Processing*, no. 11 (November).
- Bezanson, Jeff, Alan Edelman, Stefan Karpinski, and Viral B. Shah. 2017. “Julia: A Fresh Approach to Numerical Computing.” *Siam Review*, no. 1 (January).
- Bricken, Trenton, Adly Templeton, Joshua Batson, Brian Chen, Adam Jermyn, Tom Conerly, Nick Turner, et al. 2023. “Towards Monosemanticity: Decomposing Language Models with Dictionary Learning.” *Transformer Circuits Thread*.
- Higgins, I., L. Matthey, Arka Pal, Christopher P. Burgess, Xavier Glorot, M. Botvinick, S. Mohamed, and Alexander Lerchner. 2017. “Beta-VAE: Learning Basic Visual Concepts with a Constrained Variational Framework.” In *ICLR*.
- Locatello, Francesco, Stefan Bauer, Mario Lucic, Gunnar Raetsch, Sylvain Gelly, Bernhard Schölkopf, and Olivier Bachem. 2019. “Challenging Common Assumptions in the Unsupervised Learning of Disentangled Representations.” In *Proceedings of the 36th International Conference on Machine Learning*. PMLR.
- Ordonez, Vicente, Girish Kulkarni, and Tamara Berg. 2011. “Im2Text: Describing Images Using 1 Million Captioned Photographs.” In *Advances in Neural Information Processing Systems*. Curran Associates, Inc.

6 Appendix

6.1 Correlation patterns

When computing the correlations between word occurrences and dictionary assignments as described in ??, it is interesting to see if dictionary elements correlate with only few, or many words. Typically indeed we find that only

Figure 7: Word correlations of dictionary element 1.

$[-0.02, -0.01)$	3
$[-0.01, 0.0)$	1 017
$[0.0, 0.01)$	494
$[0.01, 0.02)$	7
$[0.02, 0.03)$	1
$[0.03, 0.04)$	1
$[0.04, 0.05)$	0
$[0.05, 0.06)$	0
$[0.06, 0.07)$	1
$[0.07, 0.08)$	0
$[0.08, 0.09)$	0
$[0.09, 0.1)$	1

Frequency

Figure 8: Word correlations of dictionary element 2.

$[-0.005, 0.0)$	1 265
$[0.0, 0.005)$	234
$[0.005, 0.01)$	19
$[0.01, 0.015)$	4
$[0.015, 0.02)$	2
$[0.02, 0.025)$	0
$[0.025, 0.03)$	0
$[0.03, 0.035)$	1

Frequency

Figure 9: Word correlations of dictionary element 3.

$[-0.01, -0.005)$	7
$[-0.005, 0.0)$	1 127
$[0.0, 0.005)$	343
$[0.005, 0.01)$	29

[0.01 , 0.015)	11
[0.015, 0.02)	5
[0.02 , 0.025)	1
[0.025, 0.03)	1
[0.03 , 0.035)	0
[0.035, 0.04)	1

Frequency